

Mixed Precision Block QR Implementation on GPU

University of Washington in Collaboration with Amazon Lab 126

Team

Name	Administrative Role	Technical Role	Contact
Tong Qin	Industry Mentor		tqinmath@amazon.com
Varun Elango	Faculty Mentor		varune98@uw.edu
Jaidon Lybbert	Team Lead	Researcher / CUDA	jlybbert@uw.edu
Fulin Li	Finances	Python Developer	fulin@uw.edu
Shashank Shivashankar		CUDA C++ Lead Developer	sham73@uw.edu
Alice		Python Lead Developer	alice615@uw.edu
Yiming Du	Documentation	C++ / Python Developer	yiming00@uw.edu
Mike Pao		Python Developer	mikepao@uw.edu

Background

QR factorization is the decomposition of a matrix A into the matrix product $A = QR$, where Q is an orthogonal matrix, and R is an upper triangular matrix. The factorization is used for solving nonlinear systems, and has applications in robotics, virtual and augmented reality, and 3D structure reconstruction.

The algorithm is computationally expensive, which results in high latency ill-suited for real time applications when computed sequentially. For this reason, a parallel implementation is desired, such as using CUDA on a GPU.

Mixed-precision hardware such as the Nvidia TensorCore on Volta and Turing GPU architectures, as well as Google TPUs, enable 16-bit arithmetic to reduce memory and compute workloads by sacrificing accuracy. Techniques exist to mitigate the effects of reduced accuracy while preserving the cost and energy efficiency of using mixed-precision hardware.

Objectives and Deliverables

The goal of this project is to research and implement a fast and correct mixed-precision QR factorization algorithm using mixed-precision hardware on the GPU. The measure of success is defined by the accuracy and speed of our algorithm compared to a naïve implementation on an x86 CPU architecture.

The implementation will be open source and hosted on GitHub. A stretch goal as time allows is to integrate the algorithm into an existing open-source package.

Milestones

1. Get clear idea of project scope
2. Full precision block QR implementation in Python
3. Full precision block QR implementation in CUDA
4. Prototype mixed-precision options in Python/C++
5. Fix the mixed-precision block QR algorithm
6. Fix the CUDA implementation design
7. Finish CUDA implementation
8. Test CUDA implementation
9. Final report

Timeline and Sequence of Tasks

See Appendix A. for a preliminary Gantt chart representing the project schedule.

Learning Goals

1. Develop Python skills
2. Develop CUDA C++ skills
3. Learn parallel programming concepts
4. Practice technical report writing
5. Learn to conduct and apply research

Mentor Comments

<empty>

APPENDIX A - GANTT CHART

PROJECT NAME	PROJECT LEAD	PROJECT START DATE	PROJECT END DATE	TODAY'S DATE
Mixed-precision Block QR	Jaidon Lybbert	1/3/2023	6/1/2023	1/12/2023

[illegible]